

Business Problem

Goal here is to see if we can harness the power of machine learning and boosting to help create not just a predictive model, but a general guideline for features people should look out for when picking mushrooms.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv("mushrooms.csv")
df.head()
```

Out[2]:

	class	cap- shape	cap- surface	cap- color	bruises	odor	gill- attachment	gill- spacing	gill- size	gill- color	...	stalk- surface- below- ring	...
0	p	x	s	n	t	p	f	c	n	k	...	s	
1	e	x	s	y	t	a	f	c	b	k	...	s	
2	e	b	s	w	t	l	f	c	b	n	...	s	
3	p	x	y	w	t	p	f	c	n	n	...	s	
4	e	x	s	g	f	n	f	w	b	k	...	s	

5 rows × 23 columns

The Data



Mushroom Hunting: Edible or Poisonous?

Data Source: <https://archive.ics.uci.edu/ml/datasets/Mushroom>
(<https://archive.ics.uci.edu/ml/datasets/Mushroom>)

This data set includes descriptions of hypothetical samples corresponding to 23 species of gilled mushrooms in the Agaricus and Lepiota Family (pp. 500-525). Each species is identified as definitely edible, definitely poisonous, or of unknown edibility and not recommended. This latter class was combined with the poisonous one. The Guide clearly states that there is no simple rule for determining the edibility of a mushroom; no rule like "leaflets three, let it be" for Poisonous Oak and Ivy.

Attribute Information:

1. cap-shape: bell=b,conical=c,convex=x,flat=f, knobbed=k,sunken=s
2. cap-surface: fibrous=f,grooves=g,scaly=y,smooth=s
3. cap-color: brown=n,buff=b,cinnamon=c,gray=g,green=r, pink=p,purple=u,red=e,white=w,yellow=y
4. bruises?: bruises=t,no=f
5. odor: almond=a,anise=l,creosote=c,fishy=y,foul=f, musty=m,none=n,pungent=p,spicy=s
6. gill-attachment: attached=a,descending=d,free=f,notched=n
7. gill-spacing: close=c,crowded=w,distant=d
8. gill-size: broad=b,narrow=n
9. gill-color: black=k,brown=n,buff=b,chocolate=h,gray=g, green=r,orange=o,pink=p,purple=u,red=e, white=w,yellow=y
10. stalk-shape: enlarging=e,tapering=t
11. stalk-root: bulbous=b,club=c,cup=u,equal=e, rhizomorphs=z,rooted=r,missing=?

12. stalk-surface-above-ring: fibrous=f,scaly=y,silky=k,smooth=s
13. stalk-surface-below-ring: fibrous=f,scaly=y,silky=k,smooth=s
14. stalk-color-above-ring: brown=n,buff=b,cinnamon=c,gray=g,orange=o, pink=p,red=e,white=w,yellow=y
15. stalk-color-below-ring: brown=n,buff=b,cinnamon=c,gray=g,orange=o, pink=p,red=e,white=w,yellow=y
16. veil-type: partial=p,universal=u
17. veil-color: brown=n,orange=o,white=w,yellow=y
18. ring-number: none=n,one=o,two=t
19. ring-type: cobwebby=c,evanescent=e,flaring=f,large=l, none=n,pendant=p,sheathing=s,zone=z
20. spore-print-color: black=k,brown=n,buff=b,chocolate=h,green=r, orange=o,purple=u,white=w,yellow=y
21. population: abundant=a,clustered=c,numerous=n, scattered=s,several=v,solitary=y
22. habitat: grasses=g,leaves=l,meadows=m,paths=p, urban=u,waste=w,woods=d

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   class                                8124 non-null   object
1   cap-shape                            8124 non-null   object
2   cap-surface                          8124 non-null   object
3   cap-color                            8124 non-null   object
4   bruises                             8124 non-null   object
5   odor                                8124 non-null   object
6   gill-attachment                      8124 non-null   object
7   gill-spacing                         8124 non-null   object
8   gill-size                           8124 non-null   object
9   gill-color                          8124 non-null   object
10  stalk-shape                         8124 non-null   object
11  stalk-root                          8124 non-null   object
12  stalk-surface-above-ring            8124 non-null   object
13  stalk-surface-below-ring            8124 non-null   object
14  stalk-color-above-ring              8124 non-null   object
15  stalk-color-below-ring              8124 non-null   object
16  veil-type                           8124 non-null   object
17  veil-color                          8124 non-null   object
18  ring-number                         8124 non-null   object
19  ring-type                           8124 non-null   object
20  spore-print-color                   8124 non-null   object
21  population                          8124 non-null   object
22  habitat                             8124 non-null   object
dtypes: object(23)
memory usage: 1.4+ MB
```

EDA

In []:

In []:

X & y

```
In [4]: X = pd.get_dummies(df.drop('class',axis=1),drop_first=True)
        y = df['class']
```

In [5]: X.shape

Out[5]: (8124, 95)

Train Test Split

```
In [6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15, random_state=101)
```

Modelling

Gradient Boosting with default parameters

```
In [7]: from sklearn.ensemble import GradientBoostingClassifier

gb_model = GradientBoostingClassifier()

gb_model.fit(X_train, y_train)
```

Out[7]: GradientBoostingClassifier()

prediction

```
In [8]: y_pred_train = gb_model.predict(X_train)

y_pred_test = gb_model.predict(X_test)
```

Evaluation

Accuracy


```
In [9]: from sklearn.metrics import accuracy_score
print(accuracy_score(y_train, y_pred_train)) # train accuracy
print(accuracy_score(y_test, y_pred_test))   # test accuracy

1.0
1.0
```

Confusion Matrix

```
In [10]: from sklearn.metrics import plot_confusion_matrix
print(plot_confusion_matrix(gb_model,X_test,y_test))
plt.show()
```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x000001F2CD752820>



Classification Report

```
In [11]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred_test))
```

	precision	recall	f1-score	support
e	1.00	1.00	1.00	655
p	1.00	1.00	1.00	564
accuracy			1.00	1219
macro avg	1.00	1.00	1.00	1219
weighted avg	1.00	1.00	1.00	1219

Cross Validation Score

```
In [12]: from sklearn.model_selection import cross_val_score
scores = cross_val_score(gb_model,X,y,cv=5)
print("Cross Validation Score:",scores.mean())
```

Cross Validation Score: 0.9192312239484653

Hyperparameter Tuning

```
In [13]: from sklearn.model_selection import GridSearchCV

estimator= GradientBoostingClassifier()

param_grid = {"n_estimators":[1,5,10,20,40,100], "learning_rate":[0.1,0.2,0.3,
0.5,0.8,1]}

grid = GridSearchCV(estimator, param_grid, cv=5, scoring='accuracy')

grid.fit(X_train,y_train)

grid.best_params_
```

```
Out[13]: {'learning_rate': 0.1, 'n_estimators': 100}
```

Final Model

```
In [14]: final_model = GradientBoostingClassifier(n_estimators=100,learning_rate=0.1)
final_model.fit(X_train,y_train)

preds_train = final_model.predict(X_train)
preds_test = final_model.predict(X_test)

print("Train Accuracy Score: ", accuracy_score(y_train,preds_train))
print("Test Accuracy Score: ",accuracy_score(y_test,preds_test))

Train Accuracy Score:  1.0
Test Accuracy Score:  1.0
```

Feature Importance

In [15]: `final_model.feature_importances_`

Out[15]: `array([2.91150176e-04, 4.17245225e-16, 0.00000000e+00, 0.00000000e+00, 1.22159661e-16, 1.04652037e-03, 1.17118353e-17, 3.78276239e-06, 3.48346246e-18, 1.53829642e-17, 1.13678937e-17, 0.00000000e+00, 3.73421008e-17, 1.81961056e-21, 0.00000000e+00, 5.60405971e-07, 2.31055039e-03, 5.30392238e-02, 1.84253604e-04, 1.00808510e-02, 1.82499853e-02, 3.66222069e-04, 6.14762854e-01, 9.20844491e-04, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 1.30047904e-02, 1.03950811e-02, 0.00000000e+00, 5.22781393e-17, 0.00000000e+00, 4.80274821e-17, 0.00000000e+00, 0.00000000e+00, 9.59297515e-18, 1.16274113e-16, 0.00000000e+00, 1.86446690e-17, 1.71395576e-21, 2.48611630e-18, 4.58875925e-04, 1.35972688e-01, 7.71855052e-03, 5.17747365e-02, 4.65375385e-04, 6.12113083e-06, 1.15245842e-04, 0.00000000e+00, 0.00000000e+00, 1.75588775e-02, 2.97761690e-06, 0.00000000e+00, 0.00000000e+00, 1.22682934e-03, 0.00000000e+00, 0.00000000e+00, 6.23301226e-04, 7.74443653e-05, 1.26707456e-03, 0.00000000e+00, 0.00000000e+00, 5.33104127e-05, 0.00000000e+00, 0.00000000e+00, 1.42943863e-03, 3.02342639e-03, 0.00000000e+00, 1.35380870e-07, 0.00000000e+00, 2.58827766e-03, 0.00000000e+00, 6.46948291e-05, 1.76797782e-05, 1.05423536e-05, 5.21100751e-04, 1.13072397e-02, 2.14184641e-04, 2.08997222e-04, 0.00000000e+00, 3.04953583e-02, 4.1000880e-03, 2.37566832e-03, 0.00000000e+00, 1.17406172e-03, 0.00000000e+00, 6.07840776e-08, 1.35041668e-05, 4.67493807e-04, 1.55987642e-17, 0.00000000e+00, 5.45257692e-17, 0.00000000e+00, 1.00485103e-05, 0.00000000e+00])`

In [16]: `imp_feats = pd.DataFrame(index=X.columns,data=grid.best_estimator_.feature_importances_,columns=['Importance'])`
`imp_feats`

Out[16]:

	Importance
cap-shape_c	2.911502e-04
cap-shape_f	1.772087e-16
cap-shape_k	0.000000e+00
cap-shape_s	0.000000e+00
cap-shape_x	1.021728e-16
...	...
habitat_l	0.000000e+00
habitat_m	2.673433e-17
habitat_p	0.000000e+00
habitat_u	1.004851e-05
habitat_w	0.000000e+00

95 rows × 1 columns

In [17]: `imp_feats = imp_feats[imp_feats['Importance'] > 0.01]`


```
In [18]: plt.figure(figsize=(14,6),dpi=200)
sns.barplot(data=imp_feats.sort_values('Importance'),x=imp_feats.index,y='Importance')
plt.xticks(rotation=90);
```

