

Importing Libraries

```
In [1]: import numpy as np  
import pandas as pd
```

Import Tensorflow Library

Installing Tensorflow

- pip install Tensorflow --> in anaconda prompt
- !pip install tensorflow --> in jupyter notebook

```
In [2]: import tensorflow as tf
```

```
In [3]: #tf.__version__
```

```
Out[3]: '2.4.1'
```

Keras Library

Installing Keras

- pip install Keras --> in anaconda prompt
- !pip install --upgrade keras --> in jupyter notebook

```
In [ ]: # from tensorflow import keras
```

```
In [8]: import keras
```

Part 1 - Data Preprocessing

In [9]: *### Importing the dataset*

```
dataset = pd.read_csv('Churn_Modelling.csv')
dataset
```

Out[9]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	
...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	

10000 rows × 14 columns

```
In [10]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   RowNumber        10000 non-null   int64  
 1   CustomerId      10000 non-null   int64  
 2   Surname          10000 non-null   object  
 3   CreditScore     10000 non-null   int64  
 4   Geography        10000 non-null   object  
 5   Gender           10000 non-null   object  
 6   Age              10000 non-null   int64  
 7   Tenure           10000 non-null   int64  
 8   Balance          10000 non-null   float64 
 9   NumOfProducts    10000 non-null   int64  
 10  HasCrCard       10000 non-null   int64  
 11  IsActiveMember  10000 non-null   int64  
 12  EstimatedSalary 10000 non-null   float64 
 13  Exited          10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [11]: X = dataset.iloc[:, 3:-1].values
y = dataset.iloc[:, -1].values
```

```
In [12]: X.shape, y.shape
```

```
Out[12]: ((10000, 10), (10000,))
```

```
In [13]: # Encoding categorical data
from sklearn.preprocessing import LabelEncoder
labelencoder_X_1 = LabelEncoder()
X[:, 1] = labelencoder_X_1.fit_transform(X[:, 1])

labelencoder_X_2 = LabelEncoder()
X[:, 2] = labelencoder_X_2.fit_transform(X[:, 2])
```

```
In [14]: #One Hot Encoding the "Geography" column
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [1])], remainder='passthrough')
X = np.array(ct.fit_transform(X))
```

```
In [15]: ### Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

```
In [16]: X_train.shape,X_test.shape,y_train.shape,y_train.shape
```

```
Out[16]: ((8000, 12), (2000, 12), (8000,), (8000,))
```

```
In [17]: ### Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Part 2 - Building the ANN

```
In [18]: import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
```

Initializing the ANN

```
In [19]: ann = Sequential()

#ann = keras.models.Sequential()
```

Adding the input layer and the first hidden layer

```
In [20]: ann.add(Dense(input_dim = 12,units = 6, kernel_initializer = 'uniform', activation = 'relu'))  
#ann.add(Dropout(rate = 0.1))  
  
#ann.keras.layers.add(Dense(units=6, activation='relu'))
```

Adding the second hidden layer

```
In [21]: ann.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu'))  
#ann.add(Dropout(rate = 0.1))
```

Adding the output layer

```
In [22]: ann.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
```

Part 3 - Training the ANN

Compiling the ANN

```
In [23]: ann.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```

Training the ANN on the Training set

```
In [24]: ann.fit(X_train, y_train,batch_size = 32, epochs = 100)
```

...

Part 4 - Predictions & Evaluating the model

Making the predictions

```
In [25]: y_pred = ann.predict(X_test)
y_pred = (y_pred > 0.5)
#print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)),1))
```

Evaluating the model

```
In [26]: from sklearn.metrics import confusion_matrix, accuracy_score
print("Test Accuracy:",accuracy_score(y_test, y_pred))
confusion_matrix(y_test, y_pred)
```

Test Accuracy: 0.862

```
Out[26]: array([[1531,    64],
   [ 212,  193]], dtype=int64)
```

Cross validate the model

```
In [27]: def build_cross_classifier():
    classifier = Sequential()
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu', input_dim = 12))
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu'))
    classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
    classifier.compile(optimizer = "adam", loss = 'binary_crossentropy', metrics = ['accuracy'])
    return classifier
```

```
In [28]: from keras.wrappers.scikit_learn import KerasClassifier
classifier = KerasClassifier(build_fn = build_cross_classifier, batch_size = 10, epochs = 100)
```

```
In [29]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = classifier, X = X_train, y = y_train, cv = 5)
accuracies.mean()
```

...

Predicting the result of a single observation

Use our ANN model to predict if the customer with the following information will leave the bank:

Geography: France

Credit Score: 600

Gender: Male

Age: 40 years old

Tenure: 3 years

Balance: \$ 60000

Number of Products: 2

Does this customer have a credit card ? Yes

Is this customer an Active Member: Yes

Estimated Salary: \$ 50000

So, should we say goodbye to that customer ?

Solution

```
In [30]: print(ann.predict(sc.transform([[1, 0, 0, 600, 1, 40, 3, 60000, 2, 1, 1, 50000]])) > 0.5)
[[False]]
```

Therefore, our ANN model predicts that this customer stays in the bank!

Important note 1: Notice that the values of the features were all input in a double pair of square brackets. That's because the "predict" method always expects a 2D array as the format of its inputs. And putting our values into a double pair of square brackets makes the input exactly a 2D array.

Important note 2: Notice also that the "France" country was not input as a string in the last column but as "1, 0, 0" in the first three columns. That's because of course the predict method expects the one-hot-encoded values of the state, and as we see in the first row of the matrix of features X, "France" was encoded as "1, 0, 0". And be careful to include these values in the first three columns, because the dummy variables are always created in the first columns.

Part 5 - Improving and Tuning the ANN

This can be done by using 3 options

- 1.Hyperparameter Tuning
- 2.Regularization (L1 & L2) to reduce overfitting if needed (for Regression problems)
- 3.Dropout

Hyperparameter tuning can be done for identifying no.of hidden layers & no.of neurons in each hidden layer

- layers = [[20],[10],[30],[40]] --> 1 hidden layer with different options of no.of neurons in that hidden layer
- layers = [[20],[40,20],[45,30,15]] --> Multiple hidden layers with different options of no.of neurons

Hyperparameter tuning can be done for identifying best activation function for hidden layers

- activations=['relu','sigmoid']

Hyperparameter tuning can be done for identifying best optimizers

- optimizer ='adam','rmsprop'

Hyperparameter tuning can be done for identifying best batchsize for building/training model

- batchsize =[10,16,32,64,128,256]

```
In [31]: def build_classifier(optimizer):
    classifier = Sequential()
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu', input_dim = 12))
    classifier.add(Dense(units = 6, kernel_initializer = 'uniform', activation = 'relu'))
    classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
    classifier.compile(optimizer = optimizer, loss = 'binary_crossentropy', metrics = ['accuracy'])
    return classifier
```

```
In [32]: classifier = KerasClassifier(build_fn = build_classifier)
```

```
In [33]: parameters = {'batch_size': [10,32], 'epochs': [50,100], 'optimizer': ['adam','rmsprop']}
```

```
In [34]: from sklearn.model_selection import GridSearchCV  
grid = GridSearchCV(estimator = classifier, param_grid = parameters, scoring = 'accuracy', cv = 5)
```

```
In [ ]: grid_result = grid.fit(X_train, y_train)
```

...

```
In [36]: #best parameters  
grid_result.best_params_
```

```
Out[36]: {'batch_size': 32, 'epochs': 50, 'optimizer': 'adam'}
```

```
In [37]: #best accuracy  
grid_result.best_score_
```

```
Out[37]: 0.844125
```

```
In [38]: y_pred = grid_result.predict(X_test)  
y_pred = (y_pred > 0.5)
```

```
C:\Users\rrrr90\Anaconda3\lib\site-packages\tensorflow\python\keras\engine\sequential.py:450: UserWarning: `model.predict_classes()` is deprecated and will be removed after 2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)` , if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")` , if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).  
warnings.warn(`model.predict_classes()` is deprecated and '
```

```
In [39]: # Making the Confusion Matrix  
from sklearn.metrics import confusion_matrix, accuracy_score  
  
print("Test Accuracy:",accuracy_score(y_test, y_pred))  
confusion_matrix(y_test, y_pred)
```

Test Accuracy: 0.8405

```
Out[39]: array([[1530,    65],  
                 [ 254,   151]], dtype=int64)
```